

AI in Clinical Practice.

Which tools are built for medicine, how to ground a chatbot in real evidence, and the bedside workflows worth adopting this week.

CLINICAL-GRADE TOOLS

CONNECTORS

BEDSIDE WORKFLOWS

GUARDRAILS

Dr. Abhishek J. Benur · Dr. Saikat Banerjee

Pulmonologist · Co-Founder, Parabox AI

@abhishekjbenur

PULMOCON 2026 · 6:10 PM

THE NEXT 15 MINUTES

Four questions every clinician is asking



01

Which tools are made for medicine?

Clinical-grade answer engines versus general chatbots.



02

How do I stop it making things up?

Connect ChatGPT and Claude to PubMed and Consensus.



03

Where does it help on the ward?

Differentials, discharge summaries, letters, patient education.



04

What must I never do?

Privacy, verification and accountability, in one checklist.

The promise: **every tool in this talk works on your phone tonight**, and each one comes with a guardrail.

THE LANDSCAPE

Two families of **clinical AI**



GENERAL-PURPOSE ASSISTANTS

ChatGPT · Claude · Gemini

- Brilliant at drafting, explaining, summarising and reasoning
- Trained on the open internet: answers sound right even when they're not
- Become evidence tools only when you connect them to sources



CLINICAL-GRADE ANSWER ENGINES

OpenEvidence · UpToDate Expert AI · Vera · Glass

- Built for clinicians, searching curated literature and guidelines
- Every answer cited, so you can open the source
- Verified-clinician access; some are free, some are institutional

Rule of thumb: **a clinical question goes to a clinical-grade tool first**; a writing or thinking task goes to an assistant.

CLINICAL-GRADE TOOLS

The line-up at the **point of care**

**OPENEVIDENCE****Cited clinical answers**

Searches peer-reviewed journals and guidelines; free for verified clinicians. A free version is rolling out to ~100 low- and middle-income countries (2026).

**UPTODATE EXPERT AI****Answers from UpToDate**

GenAI answers drawn only from UpToDate's expert-authored content, with Lexidrug drug information. Institutional licence; Pro Plus for individuals.

**VERA HEALTH****Evidence-graded answers**

Retrieves from 60M+ papers and guidelines, grades the evidence, cites it. Free for licensed clinicians and medical students.

**GLASS HEALTH****DDx and A&P drafting**

Keeps an evolving differential and drafts an assessment and plan with citations. Includes an ambient scribe. Free Lite tier.

Access, pricing and availability in India change often. **Check each tool's current terms before you rely on it.**

THE BIG ASSISTANTS, MADE CLINICAL

Healthcare editions of Claude and ChatGPT



ANTHROPIC · JAN 2026

Claude for Healthcare

- Native connectors: PubMed, ICD-10, CMS Coverage Database
- HIPAA-ready enterprise offering
- Agent Skills for repeatable clinical workflows



OPENAI · APR 2026

ChatGPT for Clinicians

- Free for verified US physicians, NPs, PAs and pharmacists
- Cited clinical search, documentation, reusable skills
- Not yet a default option for Indian clinicians



OPENAI · JUL 2026

ChatGPT Health

- Consumer health space inside ChatGPT (US first)
- Your patients are already using it
- Ask them what it told them

What you can use in India today: Claude or ChatGPT with the PubMed and Consensus connectors. That's next.

01

GROUNDING

Connect the chatbot to **real evidence**

A connector lets ChatGPT or Claude search PubMed or Consensus while it answers, so you get PMIDs you can open, not references it remembers.

HOW A CONNECTOR WORKS

From a memory to a **search**

01**You ask**

A focused clinical question, with the population and the outcome you care about.

›

02**The assistant plans**

It turns your question into search terms and calls the connector.

›

03**PubMed / Consensus search**

Real records come back: titles, abstracts, PMIDs, DOIs, study types.

›

04**It answers from them**

A summary that cites the retrieved papers, which you can open and check.

›

05**You verify**

Open the key papers. Check that each claim matches what the paper says.

Without a connector, the model answers from memory. **That's where invented references come from.**

SWITCH THEM ON

Set up once, use **every day**

ASSISTANT	CONNECTOR	HOW TO ENABLE	BEST FOR
Claude	PubMed	Settings → Connectors → PubMed → Connect	PMID-level searches, recent trials, full text where open access
Claude	Consensus	Settings → Connectors → Consensus → Connect	What does the evidence say? Yes/no questions across studies
ChatGPT	Consensus app	Apps (connectors) → search Consensus → Connect	Evidence summaries and study lists inside ChatGPT
ChatGPT / Claude	Deep Research	Choose Deep Research / Research in the chat box	Multi-source reports in 5–10 minutes; verify every citation

Menu names move between app versions. Look for 'Connectors' or 'Apps' in Settings.

RAPID DEMO

A clinical question, answered from PubMed

PROMPT · CLAUDE + PUBMED CONNECTOR

Search PubMed for randomised trials and meta-analyses from the last 5 years on **[high-flow nasal oxygen vs non-invasive ventilation in acute hypoxaemic respiratory failure]**.

For each study give: design, population, n, primary outcome, main result with effect size and 95% CI, and the PMID.

Then summarise what the evidence shows in 5 bullets, and say where the studies disagree.

RULES: Cite only papers you retrieved in this search. If a detail is not in the abstract, write **[NOT REPORTED]**. Do not recommend treatment for an individual patient.

WHY IT WORKS

Scoped search

Design filter and a 5-year window keep it current.

Fixed extraction

The same fields for every paper make them comparable.

Honest gaps

[NOT REPORTED] instead of an invented number.

Your judgement stays

It summarises evidence; you apply it.

02

AT THE BEDSIDE

Workflows worth adopting this week

The tasks where AI already saves time: thinking wider, writing faster, explaining better.

DIFFERENTIAL DIAGNOSIS

A Claude **Skill** for the differential

SAVED AS A CLAUDE SKILL · RUNS EVERY TIME

ROLE: Senior physician teaching diagnostic reasoning.

INPUT: De-identified case: [age, sex, presenting complaint, duration, key positives and negatives, vitals, labs, imaging].

STEPS:

1. One-line problem representation.
2. Prioritised differential: most likely, and the can't-miss diagnoses.
3. For each: findings for, findings against.
4. The next 3 tests that best separate them, and why.
5. Red flags that need action now.

RULES: Mark guideline thresholds [VERIFY]. Say what information would change the ranking.

WHY IT WORKS**What a Skill is**

Saved instructions Claude loads automatically whenever the task fits.

Same structure, every case

No re-typing the prompt on a busy ward.

Built-in safety

Can't-miss diagnoses and red flags are always listed.

Decision support only

You examine the patient and decide.

Create one: **Settings** → **Capabilities** → **Skills**, or ask Claude to 'turn this prompt into a skill'.

DOCUMENTATION

Discharge summary in **minutes, not an hour**

PROMPT · DISCHARGE SUMMARY

Act as a senior resident writing a discharge summary for a **[government teaching hospital]** in **[India]**.

MATERIAL (de-identified): **[paste admission notes, daily progress, key labs and imaging, procedures, discharge medications]**

Write it under these headings: diagnosis; presenting complaint; course in hospital; key investigations (with dates as Day 1, Day 2...); procedures; condition at discharge; discharge medications (name, dose, frequency, duration); follow-up and red-flag symptoms for the patient.

RULES: Use only the material I pasted. Write **[MISSING]** where information is absent. Do not change any drug or dose; flag apparent errors instead. Add a 5-line patient version in simple **[Hindi / Kannada / English]**.

WHY IT WORKS**No identifiers**

Day 1, Day 2 instead of dates; no name or ID number.

Grounded

'Use only the material' stops invented events.

[MISSING] flags

Gaps become visible, not papered over.

Two audiences

The chart version and the patient version together.

MORE DOMAINS

Six more places AI earns its keep



Patient education

Explain inhaler technique or a new diagnosis in the patient's language, at a 6th-grade level.



Referral and letters

Referral letters, fitness certificates, insurance and pre-authorisation drafts from your notes.



Medicines

Interaction and renal-dose checks via UpToDate Expert AI or OpenEvidence; always confirm in a formulary.



Reports and results

Summarise a long PFT, CT or lab trend into what changed and what matters.



Handover

Turn a ward list into structured SBAR handover, with pending tasks flagged.



Teaching rounds

Case-based MCQs and viva questions from today's patients, de-identified.

DECISION GUIDE

Which tool do I **actually** open?

TASK	FIRST CHOICE	WHY
A clinical question at the bedside	OpenEvidence · Vera · UpToDate Expert AI	Curated sources, cited answers, built for clinicians
What does the latest evidence say?	Claude / ChatGPT + PubMed or Consensus	Live search with PMIDs you can open
Thinking wider on a differential	Claude Skill · Glass Health	Structured, can't-miss diagnoses always listed
Discharge summary or letter	Claude · ChatGPT	Strong writers; you paste de-identified notes
Patient explanation in Hindi or a regional language	ChatGPT · Claude · Gemini	Good multilingual drafts; check the medical terms
Drug interaction or dose	UpToDate Expert AI (Lexidrug) · formulary	Drug-database backed; never a general chatbot alone

03

GUARDRAILS

Use it without getting hurt

Three failure modes and one checklist.

KNOW THE FAILURE MODES

Where clinical AI goes wrong



Confident fabrication

Plausible doses, trials and references that don't exist. Most likely when there's no connector or source.



Out of date or out of place

Old guidelines, US-centric drugs and doses, no ICMR or NTEP context unless you give it.



Automation bias

A fluent answer makes us stop thinking. Ask 'what else could this be?' before you accept it.

Counter-move for all three: **ask for sources, state your setting, and keep the final decision yours.**

NON-NEGOTIABLE

Patient data **never** goes into a public chatbot



THE RULE

De-identify before you paste

- No names, UHID or IP numbers, phone numbers, addresses or exact dates
- Crop identifiers from scans and photos
- India's DPDP Act 2023 makes health data your legal responsibility



DO THIS INSTEAD

Safe habits

- '62-year-old man, Day 3 of admission'
- Prefer enterprise or HIPAA-ready tools where your hospital provides them
- Follow your institution's AI policy; ask if none exists

BEFORE IT REACHES A PATIENT

The clinician's **verification checklist**



Sources are real

Each PMID or DOI opens, and the paper says what's claimed.



Doses are checked

Against a formulary or guideline, never the chatbot alone.



It's current

The latest guideline version, not last decade's.



It fits your setting

Indian epidemiology, drug availability, local protocols.



You understand it

If you can't explain it, you can't act on it.



A human signs off

Accountability stays with the clinician. Always.

TAKE-HOME

AI drafts. You decide.

Use clinical-grade tools for clinical questions, connect your assistant to real evidence, and keep your judgement and your patient's privacy at the centre.

UP NEXT

Difficult Cases: How AI Can Help (live cases) · Dr. Saikat Banerjee · 6:25 PM