

LIVE PROMPT-REFINEMENT EXERCISE

Ask better questions.

The same AI gives a textbook answer or a useless one. The difference is the prompt, and prompting is a skill you can learn in ten minutes.

WHAT THE LABS SAY

P-R-O-M-P-T

BAD VS GOOD

LIVE EXERCISE

Dr. Abhishek J. Benur

Pulmonologist · Co-Founder, Parabox AI · Author, The 1,000 Medical Prompts

@abhishekjbenur

PULMOCON 2026 · 7:20 PM

SAME MODEL, SAME MINUTE

The answer is only as good as **the question**

WHAT MOST OF US TYPE

Tell me about COPD exacerbation treatment.

- A generic textbook paragraph
- No patient, no setting, no severity
- US or European drugs and doses
- No sources you can check
- You still have to do all the thinking

A PROMPT WITH CONTEXT

Act as a pulmonologist in an Indian district hospital. 68-year-old man, known COPD, 3 days of increased breathlessness and purulent sputum, SpO2 86% on room air, no NIV available. Give a stepwise management plan for the first 6 hours with drugs and doses [VERIFY], when to refer, and 3 red flags. Cite GOLD 2025 where relevant.

- A plan for this patient, in this setting
- Doses flagged for you to check
- Referral triggers and red flags
- A source you can open

Nothing changed in the model. **Everything** changed in the question.

WHY PROMPTS MATTER

The AI **doesn't** know your patient



It predicts, it doesn't know

A language model writes the most likely next words. With a vague question, 'most likely' means the average textbook answer.



It fills gaps with defaults

No setting given? It assumes a well-resourced Western hospital: drugs, doses, investigations you may not have.



It sounds sure either way

A confident tone is not evidence. Unless you ask for sources and uncertainty, it won't volunteer them.

Your job is to supply **what only you know**: the patient, the setting, the purpose, and the format you need.

01

THE RULES

What **every AI lab** agrees on

Anthropic, OpenAI, Google and Microsoft publish prompting guides. They use different words for the same six ideas.

FOUR LABS, ONE CONSENSUS

Six principles from the **official guides**



Give it a role

A senior expert with a clear level and style. Google: Persona. OpenAI: Identity.



Give context and why

The patient, setting and purpose. Explaining why helps it generalise (Anthropic).



One clear task, in steps

Numbered steps beat a paragraph. Break big jobs into a chain of prompts.



Name the output

Format, length, audience and tone. Google: Format. Microsoft: Expectations.



Ground it in a source

Paste or connect the guideline or paper; ask it to quote before it answers.



Iterate

Refine in the same chat: shorter, add a case, cite that, make it a table.

SAME IDEAS, DIFFERENT LABELS

Each lab's formula

LAB	THEIR FORMULA	THE TIP THAT MATTERS MOST
Google (Gemini)	Persona · Task · Context · Format	Write in full sentences; most good prompts are ~20+ words, not 5
Microsoft (Copilot)	Goal · Context · Expectations · Source	Name the source you want it to use
OpenAI (ChatGPT)	Identity · Instructions · Examples · Context	Newer models follow instructions literally: never contradict yourself
Anthropic (Claude)	Clear, direct, with context, examples and structure	Put long documents first and your question last; use labelled sections

Our version for medicine combines all four: P-R-O-M-P-T.

THE FRAMEWORK

Every good medical prompt is **P-R-O-M-P-T**

P · PERSONA**Who the AI should be**

"A pulmonologist in a government teaching hospital who teaches residents."

R · REQUEST**One task, numbered steps**

"1. Problem representation. 2. Differential. 3. Next tests."

O · OUTPUT**Shape, length, audience**

"A table, then 5 bullets, for final-year MBBS students."

M · MATERIAL**Your context and sources**

The de-identified case, the guideline, the paper. Always before the task.

P · PARAMETERS**Rules and guardrails**

"Mark doses [VERIFY]. Say [NOT REPORTED] if it isn't in the paper."

T · TWEAK**Refine and self-check**

"Now critique your answer as a strict examiner."

ANATOMY OF A CLINICAL PROMPT

One prompt, every block labelled

PROMPT · APPROACH TO A SYMPTOM

PERSONA: You are a senior pulmonologist who teaches MD residents.

MY CONTEXT:

Setting: [tertiary government hospital, India]

Patient (de-identified): [45-year-old woman, 2 months of haemoptysis, 5 kg weight loss, non-smoker]

What I need: [a teaching-ward approach, not a final diagnosis]

TASK:

1. Confirm true haemoptysis versus mimics.
2. Prioritised differential with Indian epidemiology (TB first).
3. Stepwise investigations.
4. Red flags for massive haemoptysis.

RULES: Mark thresholds [VERIFY]. Ask me for missing details with [ASK ME].

OUTPUT: Headings, then a 60-second summary I can say on rounds.

WHY IT WORKS

Persona

Sets level, tone and specialty.

Context before task

The model reads the patient first, then the question.

Numbered task

Nothing gets skipped.

Rules with tags

Uncertainty becomes visible.

Output

Built for how you'll use it.

02

BAD VS GOOD

Three clinical **before-and-afters**

BEFORE AND AFTER · DIFFERENTIAL DIAGNOSIS

From a list of causes to a reasoning partner

B A D

Causes of fever with low platelets?

- 40 causes, alphabetical, no priorities
- Nothing about your season or region
- No 'can't-miss' flag
- No next step

G O O D

Act as an internal-medicine consultant in [coastal Karnataka, monsoon season]. 24-year-old man, 5 days of fever, myalgia, platelets 48,000, mild transaminitis, no bleeding. Give a prioritised differential (dengue, leptospirosis, scrub typhus, malaria and others), the one test that best separates each, red flags for admission, and what would change your ranking.

- Ranked for this patient, place and season
- Tests that actually discriminate
- Safety net built in

BEFORE AND AFTER · PATIENT EDUCATION

From jargon to words the patient can use

B A D

Explain asthma to a patient.

- A Wikipedia-style paragraph
- Medical terms left unexplained
- English only
- Nothing on what to actually do

G O O D

Explain asthma to a 35-year-old mother of two who studied till class 8, in simple [Hindi]. Use 6 short points: what asthma is, why the inhaler is daily, how to use a spacer, 3 warning signs to come to hospital, and one myth to correct ('inhalers are addictive'). Under 150 words. No medical jargon.

- Right language and reading level
- Behaviour, not just facts
- Addresses a real Indian myth

BEFORE AND AFTER · EVIDENCE QUESTION

From an opinion to a cited answer

B A D

Is prone positioning good?

- Good for whom? For what outcome?
- Answers from memory, may invent trials
- No effect sizes, no uncertainty

G O O D

Using the PubMed connector, find RCTs and meta-analyses on prone positioning in [awake, non-intubated adults with COVID-19 or ARDS] for the outcome [intubation]. For each: design, n, effect size with 95% CI, PMID. Then summarise in 4 bullets, noting where results conflict. Write [NOT REPORTED] for missing data.

- PICO-shaped question
- Real papers with PMIDs
- Honest about conflicting results

Turn every clinical question into **PICO**: Population, Intervention, Comparison, Outcome.

MAKE IT ADMIT UNCERTAINTY

Accuracy tags that **stop fabrication**

[VERIFY]

Numbers to check

Doses, cut-offs, statistics you must confirm in a primary source.

[CITE]

Claims needing a source

Only from a paper you uploaded or a real PMID or DOI.

[NOT REPORTED]

Honest gaps

Instead of inventing a value the paper doesn't give.

[ASK ME]

Missing information

It asks you instead of assuming.

Add one line to any prompt: "If you are unsure, say so and mark it [VERIFY]."

DON'T RESTART, REFINE

Five follow-ups that **upgrade any answer**



"Ask me 3 questions first."

It gathers context before answering.



"Critique your answer."

As a strict examiner or reviewer.



"Make it a table."

Or 5 bullets, or a one-liner.



"What did you assume?"

Hidden assumptions come out.



"Show your sources."

Then open them yourself.

The best users treat the first answer as **a draft to interrogate**, not a verdict.

LIVE EXERCISE · YOUR TURN

Refine this prompt in **three rounds**

Starting prompt: "Write about pneumonia." Shout out what to add at each round.

01**Round 1: add Persona and Material**

Who should answer?
Which patient, which setting? Community-acquired, 70-year-old, CURB-65 of 3, district hospital.

›

02**Round 2: add Request**

What exactly do we need?
Admit or not, empirical antibiotics, investigations, monitoring.

›

03**Round 3: add Parameters and Output**

Doses marked [VERIFY], cite the guideline, a one-page table for the ward, then a 3-line handover.

›

04**Compare**

Run round 0 and round 3 side by side on screen. Same model. Which would you trust on a night shift?

BEFORE YOU PRESS ENTER

The **six-second** checklist



Who should answer?

Role and level.



What does it need to know?

Patient, setting, source, de-identified.



What exactly is the task?

Numbered steps.



What should come out?

Format, length, audience.



What must it not do?

Tags: [VERIFY], [CITE], [NOT REPORTED].



How will I refine it?

Critique, sources, assumptions.

TAKE-HOME

Better questions. Better medicine.

Give the AI what only you know, ask for exactly what you need, and make it show its uncertainty. 1,000 ready-made clinical prompts: paraboxai.com.

UP NEXT

Break · then Finding Ideas & Research Questions · Dr. Saikat Banerjee · 7:40 PM